

Chapter 8.9.1 Order 3 B-Trees

Introduction to B-trees


In a binary tree, each node has one key and up to two children. A **B-tree** with order K is a tree where nodes can have up to K-1 keys and up to K children. The **order** is the maximum number of children a node can have. Ex: In a B-tree with order 4, a node can have 1, 2, or 3 keys, and up to 4 children. B-trees have the following properties:

- All keys in a B-tree must be distinct.
- All leaf nodes must be at the same level.
- An internal node with N keys must have N+1 children.
- Keys in a node are stored in sorted order from smallest to largest.
- Each key in a B-tree internal node has one left subtree and one right subtree. All left subtree keys are < that key, and all right subtree keys are > that key.

2-3-4 Tree Node Labels

- The keys in a 2-3-4 tree node are labeled as A, B, and C
- The child nodes of a 2-3-4 tree internal node are labeled as left, middle1, middle2, and right.
- If a node contains 1 key, then keys B and C, as well as children middle2 and right, are not used.
- If a node contains 2 keys, then key C, as well as the right child, are not used.
- A 2-3-4 tree node containing exactly 3 keys is said to be full, and uses all keys and children

Figure 8.9.1: 2-3-4 child subtree labels.



2-3-4 Tree Search Algorithm


- Given a key, a search algorithm returns the first node found matching that key, or returns null if a matching node is not found.
- Searching a 2-3-4 tree is a recursive process that starts with the root node.
- If the search key equals any of the keys in the node, then the node is returned.
- Otherwise, a recursive call is made on the appropriate child node.
- Which child node to use depends on the value of the search key in comparison to the node's keys.

Condition	Child node to search
key < node's A key	left
node has only 1 key or key < node's B key	middle1
node has only 2 keys or key < node's C key	middle2
none of the above	right

```

BTreeSearch(node, key) {
  if (node is not null) {
    if (node has key) {
      return node
    }
    if (key < node -->A) {
      return BTreeSearch(node -->left, key)
    }
    else if (node -->B is null || key < node -->B) {
      return BTreeSearch(node -->middle1, key)
    }
    else if (node -->C is null || key < node -->C) {
      return BTreeSearch(node -->middle2, key)
    }
    else {
      return BTreeSearch(node -->right, key)
    }
  }
  return null
}

```



```

//BTreeSplit
BTreeSplit(tree, node, nodeParent){

    if (node is not null){
        return null;
    }

    splitLeft = new BTreeNode(node-->A, node-->left, node-->middle);
    splitRight = new BTreeNode(node-->C, node-->middle2, node-->right);

    if ( nodeParent is not null){ 45 72      37
        BTreeInsertKeyWithChildren(nodeParent, node-->B, splitLeft, splitRight);
    }
    else {
        nodeParent = new BTreeNode(node-->B, splitLeft, splitRight);
        tree-->root = nodeParent;
    }
    return nodeParent;
}


```

```

BTreeSplit(tree, node, nodeParent) {
    if (node is not full) {
        return null
    }

    splitLeft = new BTreeNode(node->A, node->left, node->middle1)
    splitRight = new BTreeNode(node->C, node->middle2, node->right)
    if (nodeParent is not null) {
        BTreeInsertKeyWithChildren(nodeParent, node->B, splitLeft, splitRight)
    }
    else {
        nodeParent = new BTreeNode(node->B, splitLeft, splitRight)
        tree->root = nodeParent
    }
    return nodeParent
}

```



Since nodeParent is not null, the key 37 moves from node into nodeParent and the two newly allocated children are attached to nodeParent as well.

```

//BTreeInsertKeyWithChildren
BTreeInsertKeyWithChildren (parent, key, leftChild, rightChild){


    if( key < parent-->A){
        parent--> C = parent-->B;
        parent--> B = parent-->A;
        parent--> A = key;
        parent--> right = parent -->middle2;
        parent--> middle2 = parent-->middle1;
        parent--> middle1 = rightChild;
        parent--> left = leftChild;
    }
    else if (parent-->B is null || key < parent--> B){
        parent--> C = parent --> B;
        parent--> B = key;
        parent--> right = parent--> middle2;
        parent--> middle2 = rightChild;
        parent--> middle1 = leftChild;
    }
    else{
        parent--> C = key;
        parent--> right = rightChild;
        parent--> middle2 = leftChild;
    }
}
}

```



Inserting a key into a leaf node

- A new key is always inserted into a non-full leaf node.
- The table below describes the 4 possible cases for inserting a new key into a non-full leaf node.

Condition	Outcome
New key equals an existing key in node	No insertion takes place, and the node is not altered.
New key is < node's first key	Existing keys in node are shifted right, and the new key becomes node's first key.
Node has only 1 key or new key is < node's middle key	Node's middle key , if present, becomes last key, and new key becomes node's middle key.
None of the above	New key becomes node's last key.



38










Figure 8.9.1: 2-3-4 child subtree labels.







Insert 60




Insert 18



Insert 10



Insert 90



Merge Algorithm

- A B-Tree merge operates on a node with 1 key and increases the node's keys to 2 or 3 using either a rotation or fusion
- A node's 2 adjacent siblings are checked first during a merge, and if either has 2 or more keys, a key is transferred via a rotation
- Such a rotation increases the number of keys in the merged node from 1 to 2.
- If all adjacent siblings of the node being merged have 1 key, then fusion is used to increase the number of keys in the node from 1 to 3.
- The merge operation can be performed on any node that has 1 key and a non-null parent node with at least 2 keys

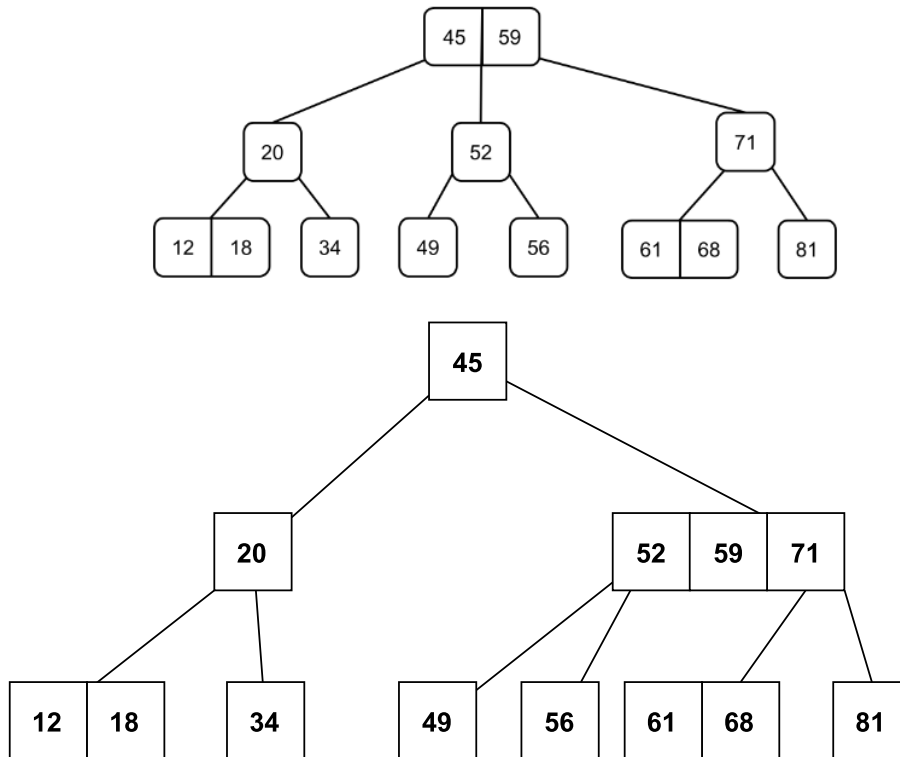




Figure 8.9.1: 2-3-4 child subtree labels.



8.13.5 BTree Remove Algorithm: non-leaf-case

BTreeRemove(tree, 60)



- When deleting 60, the process is more complex due to the key being found in an internal node
- The key 62 is a suitable replacement for 60, but 62 must be recursively removed before the swap.
- After the recursive removal completes, 60 is replaced with 62

